

Tidy Models Parameter Tuning

Setup

```
suppressPackageStartupMessages(library(tidyverse))
suppressPackageStartupMessages(library(tidymodels))
suppressPackageStartupMessages(library(ggformula))
suppressPackageStartupMessages(library(GGally))
suppressPackageStartupMessages(library(future))
library(ISLR2)

## Multitasking setup (old)
#suppressPackageStartupMessages(library(doMC))
#cores <- detectCores(logical=TRUE)
#cat('Available cores: ',cores,'\n')
#registerDoMC(cores=cores-1)

# New multitasking setup
cores <- availableCores()
cat('Available cores: ',cores,'\n')
```

Available cores: 8

```
plan(strategy=multisession, workers=cores-1)

# Load data and setup training/test/fold partitions
set.seed(2025)
hitters <- na.omit(Hitters)

hitters_split <- initial_split(hitters, strata = "Salary")
```

```
hitters_train <- training(hitters_split)
hitters_test  <- testing(hitters_split)

hitters_fold <- vfold_cv(hitters_train, v = 10)
```


[Source](#)

Exploratory analysis

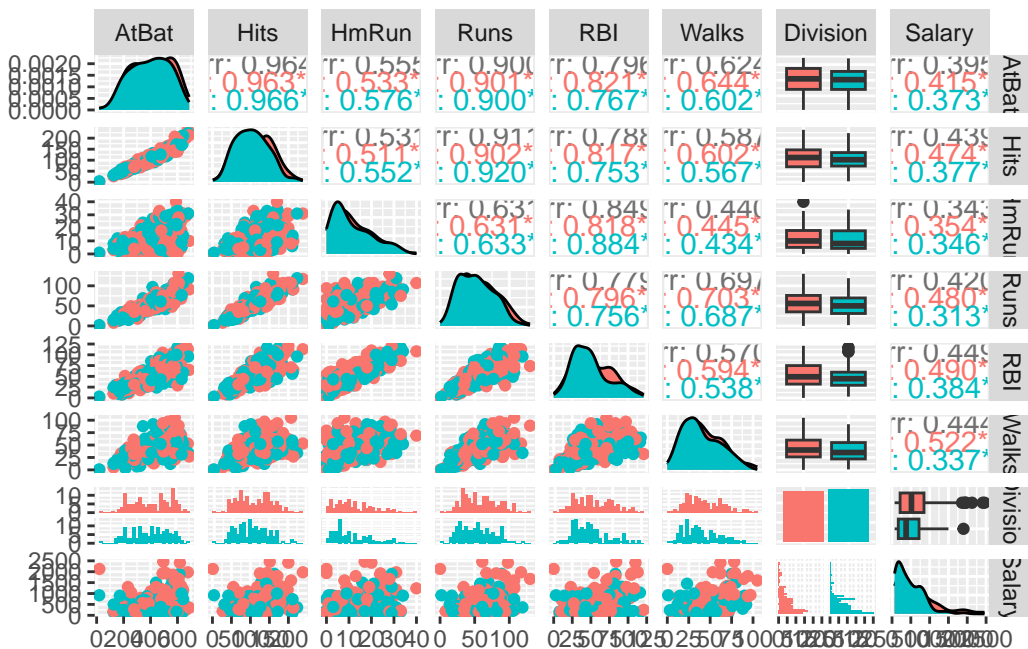
```
glimpse(hitters)
```

```
Rows: 263
Columns: 20
$ AtBat      <int> 315, 479, 496, 321, 594, 185, 298, 323, 401, 574, 202, 418, ~
$ Hits       <int> 81, 130, 141, 87, 169, 37, 73, 81, 92, 159, 53, 113, 60, 43, ~
$ HmRun      <int> 7, 18, 20, 10, 4, 1, 0, 6, 17, 21, 4, 13, 0, 7, 20, 2, 8, 16~
$ Runs       <int> 24, 66, 65, 39, 74, 23, 24, 26, 49, 107, 31, 48, 30, 29, 89, ~
$ RBI        <int> 38, 72, 78, 42, 51, 8, 24, 32, 66, 75, 26, 61, 11, 27, 75, 8~
$ Walks      <int> 39, 76, 37, 30, 35, 21, 7, 8, 65, 59, 27, 47, 22, 30, 73, 15~
$ Years      <int> 14, 3, 11, 2, 11, 2, 3, 2, 13, 10, 9, 4, 6, 13, 15, 5, 8, 1, ~
$ CAtBat     <int> 3449, 1624, 5628, 396, 4408, 214, 509, 341, 5206, 4631, 1876~
$ CHits      <int> 835, 457, 1575, 101, 1133, 42, 108, 86, 1332, 1300, 467, 392~
$ CHmRun     <int> 69, 63, 225, 12, 19, 1, 0, 6, 253, 90, 15, 41, 4, 36, 177, 5~
$ CRuns      <int> 321, 224, 828, 48, 501, 30, 41, 32, 784, 702, 192, 205, 309, ~
$ CRBI       <int> 414, 266, 838, 46, 336, 9, 37, 34, 890, 504, 186, 204, 103, ~
$ CWalks     <int> 375, 263, 354, 33, 194, 24, 12, 8, 866, 488, 161, 203, 207, ~
$ League     <fct> N, A, N, N, A, N, A, N, A, A, N, N, A, N, N, A, N, N, A, N, ~
$ Division   <fct> W, W, E, E, W, E, W, W, E, E, W, E, E, E, W, W, W, E, W, W, ~
$ PutOuts    <int> 632, 880, 200, 805, 282, 76, 121, 143, 0, 238, 304, 211, 121~
$ Assists    <int> 43, 82, 11, 40, 421, 127, 283, 290, 0, 445, 45, 11, 151, 45, ~
$ Errors     <int> 10, 14, 3, 4, 25, 7, 9, 19, 0, 22, 11, 7, 6, 8, 10, 16, 2, 5~
$ Salary     <dbl> 475.000, 480.000, 500.000, 91.500, 750.000, 70.000, 100.000, ~
$ NewLeague  <fct> N, A, N, N, A, A, A, N, A, A, N, N, A, N, N, A, N, N, N, N, ~
```

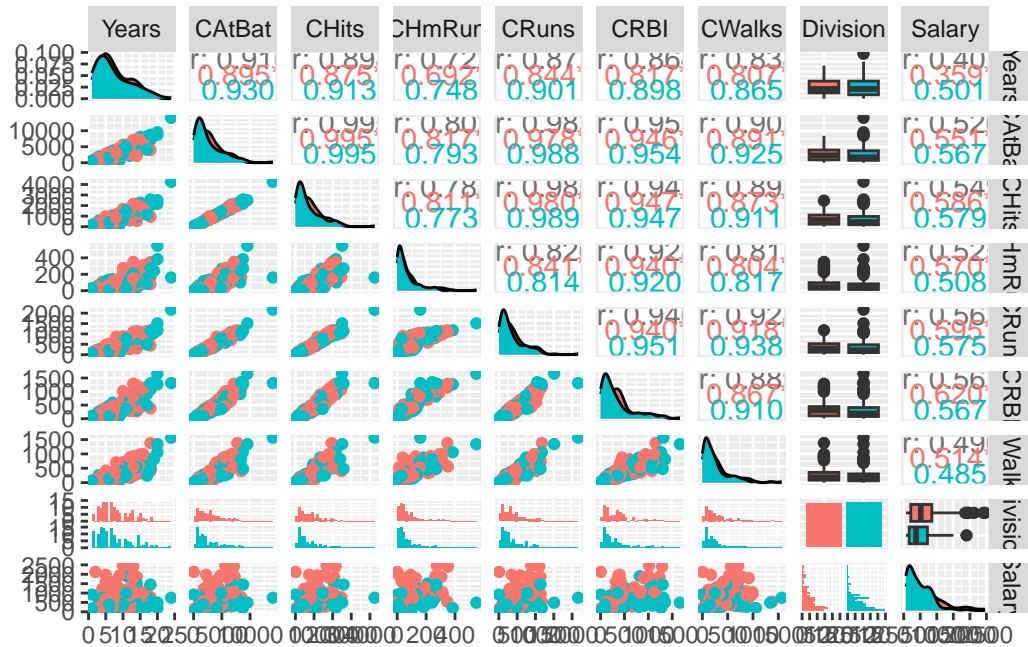
```
library(corrplot)
hitters %>%
  dplyr::select(-c(League, Division, NewLeague)) %>%
  cor() %>% corrplot()
```



```
hitters %>%
  select(AtBat:Walks, Division, Salary) %>%
  ggpairs(mapping=aes(col=Division))
```



```
hitters %>%
  select(Years:CWalks,Division,Salary) %>%
  ggpairs(mapping=aes(col=Division))
```



```
library(car)
vif(lm(Salary~.,data=hitters))
```

AtBat	Hits	HmRun	Runs	RBI	Walks	Years
22.944366	30.281255	7.758668	15.246418	11.921715	4.148712	9.313280
CAtBat	CHits	CHmRun	CRuns	CRBI	CWalks	League
251.561160	502.954289	46.488462	162.520810	131.965858	19.744105	4.134115
Division	PutOuts	Assists	Errors	NewLeague		
1.075398	1.236317	2.709341	2.214543	4.099063		

Ridge regression

```
ridge_spec <- linear_reg(mixture = 0, penalty = 0) %>%
  set_mode("regression") %>%
  set_engine("glmnet")
```

```
ridge_fit <- fit(ridge_spec, Salary ~ ., data = hitters)
tidy(ridge_fit)
```

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyr':

expand, pack, unpack

Loaded glmnet 4.1-8

A tibble: 20 x 3

	term <chr>	estimate <dbl>	penalty <dbl>
1	(Intercept)	81.1	0
2	AtBat	-0.682	0
3	Hits	2.77	0
4	HmRun	-1.37	0
5	Runs	1.01	0
6	RBI	0.713	0
7	Walks	3.38	0
8	Years	-9.07	0
9	CAtBat	-0.00120	0
10	CHits	0.136	0
11	CHmRun	0.698	0
12	CRuns	0.296	0
13	CRBI	0.257	0
14	CWalks	-0.279	0
15	LeagueN	53.2	0
16	DivisionW	-123.	0
17	PutOuts	0.264	0
18	Assists	0.170	0
19	Errors	-3.69	0
20	NewLeagueN	-18.1	0

```
tidy(ridge_fit, penalty = 11498)
```

A tibble: 20 x 3

	term	estimate	penalty
--	------	----------	---------

	<chr>	<dbl>	<dbl>
1	(Intercept)	407.	11498
2	AtBat	0.0370	11498
3	Hits	0.138	11498
4	HmRun	0.525	11498
5	Runs	0.231	11498
6	RBI	0.240	11498
7	Walks	0.290	11498
8	Years	1.11	11498
9	CAtBat	0.00314	11498
10	CHits	0.0117	11498
11	CHmRun	0.0876	11498
12	CRuns	0.0234	11498
13	CRBI	0.0242	11498
14	CWalks	0.0250	11498
15	LeagueN	0.0866	11498
16	DivisionW	-6.23	11498
17	PutOuts	0.0165	11498
18	Assists	0.00262	11498
19	Errors	-0.0206	11498
20	NewLeagueN	0.303	11498

```
tidy(ridge_fit, penalty = 705)
```

```
# A tibble: 20 x 3
```

	term	estimate	penalty
	<chr>	<dbl>	<dbl>
1	(Intercept)	54.4	705
2	AtBat	0.112	705
3	Hits	0.656	705
4	HmRun	1.18	705
5	Runs	0.937	705
6	RBI	0.847	705
7	Walks	1.32	705
8	Years	2.58	705
9	CAtBat	0.0108	705
10	CHits	0.0468	705
11	CHmRun	0.338	705
12	CRuns	0.0937	705
13	CRBI	0.0979	705
14	CWalks	0.0718	705
15	LeagueN	13.7	705


16	DivisionW	-54.7	705
17	PutOuts	0.119	705
18	Assists	0.0161	705
19	Errors	-0.704	705
20	NewLeagueN	8.61	705

```
tidy(ridge_fit, penalty = 50)
```

```
# A tibble: 20 x 3
```

	term <chr>	estimate <dbl>	penalty <dbl>
1	(Intercept)	48.2	50
2	AtBat	-0.354	50
3	Hits	1.95	50
4	HmRun	-1.29	50
5	Runs	1.16	50
6	RBI	0.809	50
7	Walks	2.71	50
8	Years	-6.20	50
9	CAtBat	0.00609	50
10	CHits	0.107	50
11	CHmRun	0.629	50
12	CRuns	0.217	50
13	CRBI	0.215	50
14	CWalks	-0.149	50
15	LeagueN	45.9	50
16	DivisionW	-118.	50
17	PutOuts	0.250	50
18	Assists	0.121	50
19	Errors	-3.28	50
20	NewLeagueN	-9.42	50

```
ridge_fit %>%  
autoplot()
```



```
predict(ridge_fit, new_data = hitters)
```

```
# A tibble: 263 x 1
  .pred
  <dbl>
1  442.
2  676.
3 1059.
4  521.
5  543.
6  218.
7   74.7
8   96.1
9  809.
10 865.
# i 253 more rows
```

```
predict(ridge_fit, new_data = hitters, penalty = 500)
```

```
# A tibble: 263 x 1
  .pred
```



```

      <dbl>
1  525.
2  620.
3  895.
4  425.
5  589.
6  179.
7  147.
8  187.
9  841.
10 840.
# i 253 more rows

```

```

ridge_recipe <-
  recipe(formula = Salary ~ ., data = hitters_train) %>%
  step_novel(all_nominal_predictors()) %>% # add 'new' level to factors and chr -> fctr
  step_dummy(all_nominal_predictors()) %>% # factors to dummy variables
  step_zv(all_predictors()) %>%           # remove zero variance predictors
  step_normalize(all_predictors())        # normalize all predictors

ridge_spec <-
  linear_reg(penalty = tune(), mixture = 0) %>%
  set_mode("regression") %>%
  set_engine("glmnet")

ridge_workflow <- workflow() %>%
  add_recipe(ridge_recipe) %>%
  add_model(ridge_spec)

penalty_grid <- grid_regular(penalty(range = c(-5, 5)), levels = 50)
penalty_grid

```

```

# A tibble: 50 x 1
  penalty
  <dbl>
1 0.00001
2 0.0000160
3 0.0000256
4 0.0000409
5 0.0000655
6 0.000105
7 0.000168

```

```

8 0.000268
9 0.000429
10 0.000687
# i 40 more rows

```

```

tune_res <- tune_grid(
  ridge_workflow,
  resamples = hitters_fold,
  grid = penalty_grid,
  control(parallel_over = "everything")
)

```

Warning: The `...` are not used in this function but one or more objects were passed: ''


```
tune_res
```

```

# Tuning results
# 10-fold cross-validation
# A tibble: 10 x 4
  splits          id    .metrics          .notes
  <list>         <chr>  <list>          <list>
1 <split [176/20]> Fold01 <tibble [100 x 5]> <tibble [0 x 3]>
2 <split [176/20]> Fold02 <tibble [100 x 5]> <tibble [0 x 3]>
3 <split [176/20]> Fold03 <tibble [100 x 5]> <tibble [0 x 3]>
4 <split [176/20]> Fold04 <tibble [100 x 5]> <tibble [0 x 3]>
5 <split [176/20]> Fold05 <tibble [100 x 5]> <tibble [0 x 3]>
6 <split [176/20]> Fold06 <tibble [100 x 5]> <tibble [0 x 3]>
7 <split [177/19]> Fold07 <tibble [100 x 5]> <tibble [0 x 3]>
8 <split [177/19]> Fold08 <tibble [100 x 5]> <tibble [0 x 3]>
9 <split [177/19]> Fold09 <tibble [100 x 5]> <tibble [0 x 3]>
10 <split [177/19]> Fold10 <tibble [100 x 5]> <tibble [0 x 3]>

```

```
autoplot(tune_res)
```



```
collect_metrics(tune_res) %>% filter(.metric=="rmse")
```

```
# A tibble: 50 x 7
```

	penalty	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	0.00001	rmse	standard	361.	10	28.3	Preprocessor1_Model01
2	0.0000160	rmse	standard	361.	10	28.3	Preprocessor1_Model02
3	0.0000256	rmse	standard	361.	10	28.3	Preprocessor1_Model03
4	0.0000409	rmse	standard	361.	10	28.3	Preprocessor1_Model04
5	0.0000655	rmse	standard	361.	10	28.3	Preprocessor1_Model05
6	0.000105	rmse	standard	361.	10	28.3	Preprocessor1_Model06
7	0.000168	rmse	standard	361.	10	28.3	Preprocessor1_Model07
8	0.000268	rmse	standard	361.	10	28.3	Preprocessor1_Model08
9	0.000429	rmse	standard	361.	10	28.3	Preprocessor1_Model09
10	0.000687	rmse	standard	361.	10	28.3	Preprocessor1_Model10

```
# i 40 more rows
```

```
best_penalty <- select_best(tune_res, metric = "rmse")
best_penalty
```

```
# A tibble: 1 x 2
```

```

penalty .config
  <dbl> <chr>
1    356. Preprocessor1_Model38

```

```

ridge_final <- finalize_workflow(ridge_workflow, best_penalty)
ridge_final_fit <- fit(ridge_final, data = hitters_train)
ridge_aug <- augment(ridge_final_fit, new_data = hitters_test)
ridge_result <- rbind(
  rmse(ridge_aug, truth = Salary, estimate = .pred),
  rsq(ridge_aug, truth = Salary, estimate = .pred))
ridge_result

```

```

# A tibble: 2 x 3
  .metric .estimator .estimate
  <chr>    <chr>        <dbl>
1 rmse    standard      277.
2 rsq     standard       0.515

```

Final ridge model

```

ridge_final_fit <- fit(ridge_final, data=hitters)
tidy(ridge_final_fit, conf_int=TRUE)

```


```

# A tibble: 20 x 3
  term          estimate penalty
  <chr>         <dbl>    <dbl>
1 (Intercept)   536.      356.
2 AtBat         12.3      356.
3 Hits         37.6      356.
4 HmRun         6.00      356.
5 Runs         26.9      356.
6 RBI          22.8      356.
7 Walks        34.5      356.
8 Years         7.44      356.
9 CAtBat       25.9      356.
10 CHits       36.4      356.
11 CHmRun      32.7      356.
12 CRuns      37.1      356.
13 CRBI       38.2      356.
14 CWalks     14.9      356.
15 PutOuts    45.0      356.


```

16 Assists	3.96	356.
17 Errors	-8.43	356.
18 League_N	10.6	356.
19 Division_W	-38.2	356.
20 NewLeague_N	4.65	356.

```
ridge_aug <- augment(ridge_final_fit, new_data=hitters) %>% mutate(.resid=Salary-.pred)
ridge_aug %>%
  gf_qq(~.resid) %>% gf_qqline()
```



```
ridge_aug %>%
  gf_point(~.resid~.pred)
```



The Lasso

```
lasso_recipe <-
  recipe(formula = Salary ~ ., data = hitters_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_normalize(all_predictors())

lasso_spec <-
  linear_reg(penalty = tune(), mixture = 1) %>%
  set_mode("regression") %>%
  set_engine("glmnet")


lasso_workflow <- workflow() %>%
  add_recipe(lasso_recipe) %>%
  add_model(lasso_spec)

penalty_grid <- grid_regular(penalty(range = c(-2, 2)), levels = 50)
```

```
tune_res <- tune_grid(
  lasso_workflow,
  resamples = hitters_fold,
  grid = penalty_grid,
  control(parallel_over = "everything")
)
```

Warning: The `...` are not used in this function but one or more objects were passed: ''

```
autoplot(tune_res)
```



```
best_penalty <- select_best(tune_res, metric = "rsq")
lasso_final <- finalize_workflow(lasso_workflow, best_penalty)
lasso_final_fit <- fit(lasso_final, data = hitters_train)
lasso_aug <- augment(lasso_final_fit, new_data = hitters_test)
lasso_result <- rbind(
  rmse(lasso_aug, truth = Salary, estimate = .pred),
  rsq(lasso_aug, truth = Salary, estimate = .pred))
lasso_result
```


```
# A tibble: 2 x 3
  .metric .estimator .estimate
  <chr>    <chr>         <dbl>
1 rmse     standard       277.
2 rsq      standard        0.493
```

Final lasso model


```
lasso_final_fit <- fit(lasso_final, data=hitters)
tidy(lasso_final_fit, conf_int=TRUE)
```

```
# A tibble: 20 x 3
  term          estimate penalty
  <chr>         <dbl>    <dbl>
1 (Intercept)   536.      3.39
2 AtBat        -198.      3.39
3 Hits         235.      3.39
4 HmRun          0       3.39
5 Runs          0       3.39
6 RBI           0       3.39
7 Walks        94.9     3.39
8 Years       -40.7     3.39
9 CAtBat         0       3.39
10 CHits         0       3.39
11 CHmRun       35.2     3.39
12 CRuns       197.      3.39
13 CRBI        130.      3.39
14 CWalks     -118.      3.39
15 PutOuts     74.7     3.39
16 Assists     19.6     3.39
17 Errors     -11.5     3.39
18 League_N     15.4     3.39
19 Division_W  -59.7     3.39
20 NewLeague_N  0         3.39
```

```
lasso_aug <- augment(lasso_final_fit, new_data=hitters) %>% mutate(.resid=Salary-.pred)
lasso_aug %>%
  gf_qq(~.resid) %>% gf_qqline()
```

```
lasso_aug %>%
  gf_point(.resid~.pred)
```



Principal components regression

```
lm_spec <-  
  linear_reg() %>%  
  set_mode("regression") %>%  
  set_engine("lm")  
  
pca_recipe <-  
  recipe(formula = Salary ~ ., data = hitters_train) %>%  
  step_novel(all_nominal_predictors()) %>%  
  step_dummy(all_nominal_predictors()) %>%  
  step_zv(all_predictors()) %>%  
  step_normalize(all_predictors()) %>%  
  step_pca(all_predictors(), threshold = tune())  
  
pca_workflow <-  
  workflow() %>%  
  add_recipe(pca_recipe) %>%  
  add_model(lm_spec)  
  
threshold_grid <- grid_regular(threshold(), levels = 10)  
threshold_grid
```

```
# A tibble: 10 x 1  
  threshold  
    <dbl>  
1         0  
2      0.111  
3      0.222  
4      0.333  
5      0.444  
6      0.556  
7      0.667  
8      0.778  
9      0.889  
10       1
```

```
tune_res <- tune_grid(  
  pca_workflow,  
  resamples = hitters_fold,
```

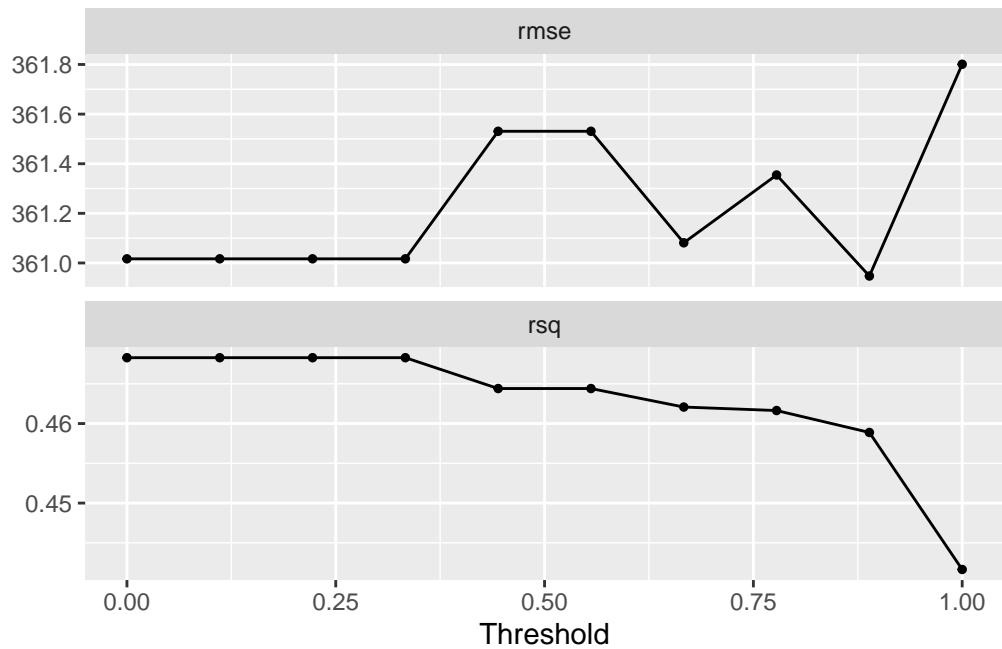
```

grid = threshold_grid,
control(parallel_over = "everything")
)

```

Warning: The `...` are not used in this function but one or more objects were passed: ''

```
autoplot(tune_res)
```



```

best_threshold <- select_best(tune_res, metric = "rmse")
pca_final <- finalize_workflow(pca_workflow, best_threshold)
pca_final_fit <- fit(pca_final, data = hitters_train)
pca_aug <- augment(pca_final_fit, new_data = hitters_test)
pca_result <- rbind(
  rmse(pca_aug, truth = Salary, estimate = .pred),
  rsq(pca_aug, truth = Salary, estimate = .pred))
pca_result

```

```

# A tibble: 2 x 3
  .metric .estimator .estimate
  <chr>   <chr>       <dbl>

```

```
1 rmse      standard      271.
2 rsq       standard      0.521
```


Final pca model

```
pca_final_fit <- fit(pca_final, data=hitters)
tidy(pca_final_fit, conf_int=TRUE)
```

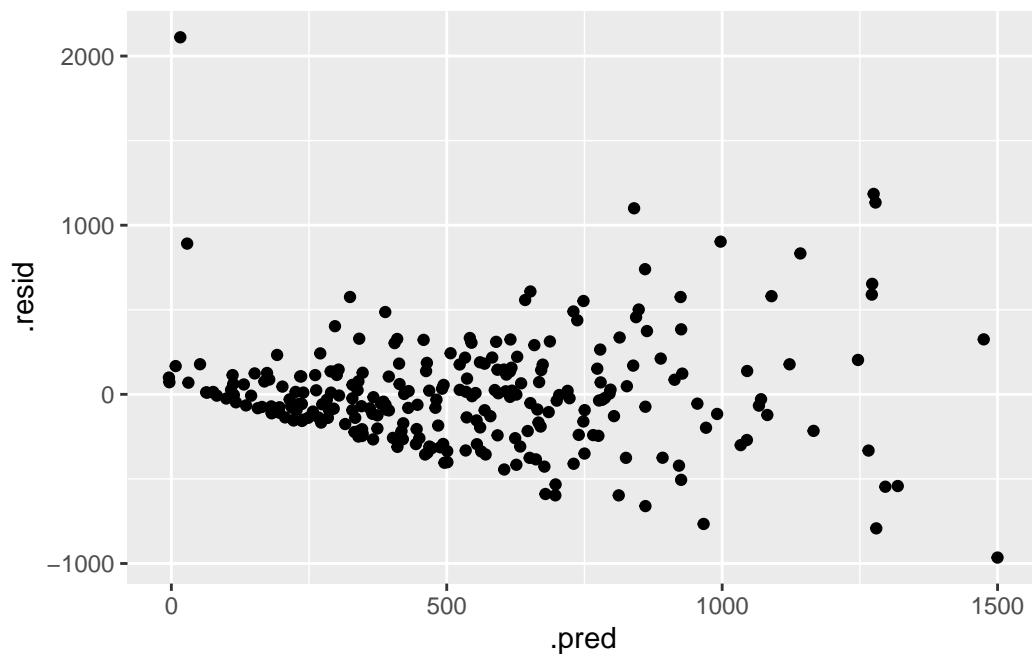
A tibble: 8 x 5

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	536.	20.6	26.0	9.20e-74
2	PC1	107.	7.65	13.9	3.24e-33
3	PC2	21.6	10.1	2.14	3.35e- 2
4	PC3	-24.3	14.5	-1.68	9.39e- 2
5	PC4	-37.1	16.5	-2.24	2.59e- 2
6	PC5	58.5	20.6	2.84	4.94e- 3
7	PC6	62.3	22.7	2.75	6.48e- 3
8	PC7	-24.7	24.8	-0.994	3.21e- 1

```
pca_aug <- augment(pca_final_fit, new_data=hitters) %>% mutate(.resid=Salary-.pred)
pca_aug %>%
  gf_qq(~.resid) %>% gf_qqline()
```



```
pca_aug %>%
  gf_point(.resid~.pred)
```



Partial least squares

Partial least squares in Tidy Models requires the package ‘mixOmics’ which is only available with Bioconductor which is a huge installation.

```
if (!require("BiocManager", quietly = TRUE))
  install.packages("BiocManager")
```

```
BiocManager::install("mixOmics")
```

```
pls_recipe <-
  recipe(formula = Salary ~ ., data = hitters_train) %>%
    step_novel(all_nominal_predictors()) %>%
    step_dummy(all_nominal_predictors()) %>%
    step_zv(all_predictors()) %>%
    step_normalize(all_predictors()) %>%
    step_pls(all_predictors(), num_comp = tune(), outcome = "Salary")
```

```

lm_spec <- linear_reg() %>%
  set_mode("regression") %>%
  set_engine("lm")

pls_workflow <- workflow() %>%
  add_recipe(pls_recipe) %>%
  add_model(lm_spec)

num_comp_grid <- grid_regular(num_comp(c(1, 20)), levels = 10)

tune_res <- tune_grid(
  pls_workflow,
  resamples = hitters_fold,
  grid = num_comp_grid,
  control(parallel_over = "everything")
)

```

Warning: The `...` are not used in this function but one or more objects were passed: ''

```

best_threshold <- select_best(tune_res, metric = "rmse")

pls_final <- finalize_workflow(pls_workflow, best_threshold)

pls_final_fit <- fit(pls_final, data = hitters_train)

pls_aug <- augment(pls_final_fit, new_data = hitters_test)
pls_result <- rbind(
  rmse(pls_aug, truth = Salary, estimate = .pred),
  rsq(pls_aug, truth = Salary, estimate = .pred))
pls_result

```

```

# A tibble: 2 x 3
  .metric .estimator .estimate
  <chr>    <chr>         <dbl>
1 rmse    standard         291.
2 rsq     standard          0.447

```


Final pls model

```
pls_final_fit <- fit(pls_final, data=hitters)
tidy(pls_final_fit, conf_int=TRUE)
```


A tibble: 12 x 5

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	536.	19.3	27.8	1.24e-78
2	PLS01	111.	7.24	15.3	6.78e-38
3	PLS02	65.6	15.3	4.28	2.69e-5
4	PLS03	34.7	12.9	2.68	7.76e-3
5	PLS04	49.5	21.3	2.32	2.10e-2
6	PLS05	99.6	31.6	3.15	1.81e-3
7	PLS06	52.5	21.1	2.49	1.36e-2
8	PLS07	59.5	30.9	1.92	5.57e-2
9	PLS08	84.7	37.8	2.24	2.61e-2
10	PLS09	37.6	31.2	1.21	2.29e-1
11	PLS10	57.2	48.8	1.17	2.42e-1
12	PLS11	48.6	40.5	1.20	2.31e-1

```
pls_aug <- augment(pls_final_fit, new_data=hitters) %>% mutate(.resid=Salary-.pred)
pls_aug %>%
  gf_qq(~.resid) %>% gf_qqline()
```



```
pls_aug %>%
  gf_point(.resid~.pred)
```



Blended Ridge and Lasso

```
mixed_recipe <-
  recipe(formula = Salary ~ ., data = hitters_train) %>%
  step_novel(all_nominal_predictors()) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_predictors()) %>%
  step_normalize(all_predictors())
prep(mixed_recipe) #>% bake(hitters_train)
```

-- Recipe -----

-- Inputs

Number of variables by role

outcome: 1
predictor: 19

-- Training information

Training data contained 196 data points and no incomplete rows.

-- Operations

* Novel factor level assignment for: League, Division, NewLeague | Trained

* Dummy variables from: League, Division, NewLeague | Trained

* Zero variance filter removed: League_new and Division_new, ... | Trained

* Centering and scaling for: AtBat, Hits, HmRun, Runs, RBI, ... | Trained

```
mixed_spec <-  
  linear_reg(penalty = tune("penalty"), mixture = tune("mixture")) %>%  
  set_mode("regression") %>%  
  set_engine("glmnet")  
  
extract_parameter_set_dials(mixed_spec)
```

Collection of 2 parameters for tuning

identifier	type	object
penalty	penalty	nparam[+]
mixture	mixture	nparam[+]

```

mixed_workflow <- workflow() %>%
  add_recipe(mixed_recipe) %>%
  add_model(mixed_spec)

mixed_grid <- grid_regular(penalty(range = c(-5, 5)), mixture(), levels = 50)
mixed_grid %>%
  pull(penalty) %>% range()

```

```
[1] 1e-05 1e+05
```

```

grid_regular(penalty(), mixture(), levels = 50) %>%
  pull(penalty) %>% range()

```

```
[1] 1e-10 1e+00
```

```

expand_grid("penalty"=seq(1e-5, 1e5, 2000), "mixture"=seq(0, 1, 1/50)) %>%
  pull(penalty) %>% range()

```

```
[1] 1.0e-05 9.8e+04
```

```

system.time(
  suppressMessages(
    tune_res <- tune_grid(
      mixed_workflow,
      resamples = hitters_fold,
      grid = mixed_grid,
      control(parallel_over = "everything")
    )
  )
)

```

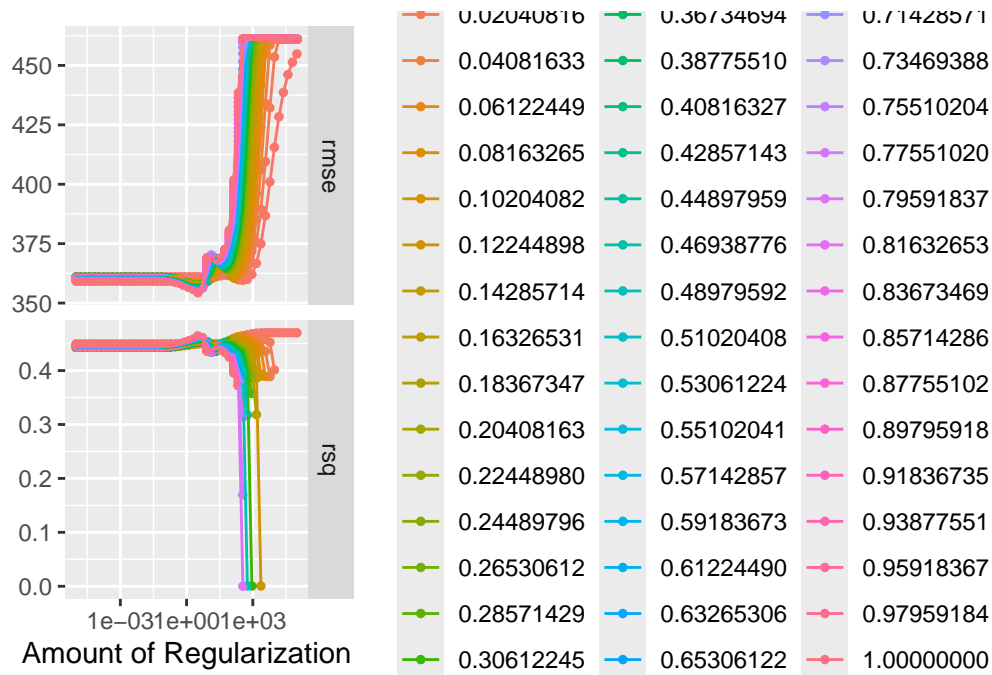
Warning: The `...` are not used in this function but one or more objects were passed: ''

```

  user  system elapsed
1.769   0.042  22.779

```

```
autoplot(tune_res)
```



```
best_penalty <- select_best(tune_res, metric = "rmse")
mixed_final <- finalize_workflow(mixed_workflow, best_penalty)
mixed_final_fit <- fit(mixed_final, data = hitters_train)
mixed_aug <- augment(mixed_final_fit, new_data = hitters_test)
mixed_result <- rbind(
  rmse(mixed_aug, truth = Salary, estimate = .pred),
  rsq(mixed_aug, truth = Salary, estimate = .pred))
mixed_result
```

```
# A tibble: 2 x 3
  .metric .estimator .estimate
  <chr>   <chr>       <dbl>
1 rmse    standard      277.
2 rsq     standard       0.491
```


Final mixed model

```
mixed_final_fit <- fit(mixed_final, data=hitters)
tidy(mixed_final_fit, conf_int=TRUE)
```


```
# A tibble: 20 x 3
```

	term <chr>	estimate <dbl>	penalty <dbl>
1	(Intercept)	536.	3.24
2	AtBat	-205.	3.24
3	Hits	240.	3.24
4	HmRun	0	3.24
5	Runs	0	3.24
6	RBI	0	3.24
7	Walks	96.6	3.24
8	Years	-41.8	3.24
9	CAtBat	0	3.24
10	CHits	0	3.24
11	CHmRun	36.6	3.24
12	CRuns	201.	3.24
13	CRBI	129.	3.24
14	CWalks	-123.	3.24
15	PutOuts	75.1	3.24
16	Assists	20.8	3.24
17	Errors	-11.9	3.24
18	League_N	15.6	3.24
19	Division_W	-59.7	3.24
20	NewLeague_N	0	3.24

```
mixed_aug <- augment(mixed_final_fit, new_data=hitters) %>% mutate(.resid=Salary-.pred)
mixed_aug %>%
  gf_qq(~.resid) %>% gf_qqline()
```



```
mixed_aug %>%  
  gf_point(.resid~.pred)
```



Results comparison

Compare models fit on the training data and assessed on the test data. The base linear model fit wasn't done above, so do it here.

```
lm_final_fit <- fit(lm_spec, Salary ~ ., data = hitters_train)
lm_aug <- augment(lm_final_fit, new_data = hitters_test)
lm_result <- rbind(
  rmse(lm_aug, truth = Salary, estimate = .pred),
  rsq(lm_aug, truth = Salary, estimate = .pred))
lm_result
```

```
# A tibble: 2 x 3
  .metric .estimator .estimate
  <chr>    <chr>         <dbl>
1 rmse    standard        298.
2 rsq     standard         0.421
```

```
results <-
  rbind(
    lm_result %>% mutate(method='lm'),
    ridge_result %>% mutate(method='ridge'),
    lasso_result %>% mutate(method='lasso'),
    pca_result %>% mutate(method='pca'),
    pls_result %>% mutate(method='pls'),
    mixed_result %>% mutate(method='mixed')) %>%
  dplyr::select(method, .metric, .estimate) %>%
  arrange(.metric, .estimate)
results
```

```
# A tibble: 12 x 3
  method .metric .estimate
  <chr>    <chr>         <dbl>
1 pca     rmse          271.
2 ridge   rmse          277.
3 lasso   rmse          277.
4 mixed   rmse          277.
5 pls     rmse          291.
6 lm      rmse          298.
7 lm      rsq            0.421
8 pls     rsq            0.447
```

9	mixed	rsq	0.491
10	lasso	rsq	0.493
11	ridge	rsq	0.515
12	pca	rsq	0.521

```
results %>%
  filter(.metric=='rmse') %>%
  gf_point(.estimate~method) %>%
  gf_labs(y='rmse',title="Cross-validated rmse by method")
```

